

Feedback for 1804170

Computer Vision Assignment

Presentation: 25/25

The appearance of your code was generally fine. The introductory comments at the top of your program were good. The comments interspersed with your code made a good job of explaining how the code works and what individual sections of it are meant to do. In fact, this is the best-commented program I have seen (so far) for the assignment — well done!

Algorithms: 22/25

You identify the various targets by colour, which is the most sensible way. You use OpenCV's `inRange` function in HSV space, a sensible choice, with some per-pixel processing to identify object limits. This approach is fine here but wouldn't work if there were to be more than one object of a particular colour. Your code for determining the distance to the meteors looks sensible.

You were able to track the motion of objects across the frames, the trickiest part of the assignment. You are using the lengths of normalized motion vectors computed using the inner product. That is close to the approach I used in my own solution, the angles between vectors computed by the inner product. My view is that this is the most general approach with the information we have been given by that lowly technician.

Coding style: 22/25

You have nicely separated some of the code out into separate routines, which makes the program as a whole easier to read, though in a few places I think the functionality could have been made more straightforward. Also nice to see some doctests and asserts in your code, something I use quite often. Incidentally, you can use

```
im = cv2.imread ("...")  
if im is None:
```

to detect `imread`'s failure to read an image; note the use of `is`.

Results: 21/25

Your program works, correctly distinguishing the alien spacecraft from the meteors — though you have identified meteors as alien spacecraft too. Fortunately, there is more than one missile to fire and you have saved the human race from what seemed like certain destruction. The entire 10 billion people on Earth breathe a collective sigh of relief and you are awarded the Galactic Order of Merit for your achievements ; -). I read with some interest your remarks about the relative danger of false positives and false negatives.

Your submission works well with the test harness. Your program also runs with my unseen test sequence, though again it identifies most of the objects as UFOs.

Although not an assessment criterion, I was impressed by how quickly your program ran.

Overall mark: 90%

```

1  """
2  ---
3  **ce316ass.py**
4  ---
5
6  ---
7  **PURPOSE:**
8  ---
9
10 Given:
11
12 * A number of frames to look at
13   * Must be at least 1. Can't really look at less than 1 frame, y'know?
14 * Templates for filenames for left/right images
15   * such as left-%03d.png and right-%03d.png
16 * Image files, which can be opened, with filenames that conform to the
17   aforementioned templates
18   * They must be of the same dimensions as each other.
19     * If the dimensions differ, the program will terminate with an error
20       message.
21   * They must be openable.
22     * If they cannot be opened, the program will terminate with an error
23       message.
24
25 This will:
26
27 * Identify objects identified in the images
28   * Assumptions
29     * Assumes that the objects have these HSV values
30     * Cyan
31       * Hue 180
32     * Red
33       * Hue 0
34       * Saturation of at least 1/256
35     * White
36       * Hue 0
37       * Saturation 0
38       * Value of at least 1/256 (and not in 'cyan')
39     * Blue
40       * Hue 240
41       * Saturation of at least 1/256
42       * Value of at least 1/256
43     * Green
44       * Hue 120
45       * Saturation of at least 1/256
46       * Value of at least 1/256
47     * Yellow
48       * Hue 60 (and not in 'orange')
49       * Saturation of at least 1/256
50     * Orange
51       * Hue 40-58
52     * Assumes there will be only one object of the given colour in the image
53     * or 0 objects of that given colour :shrug:
54     * Assumes that the objects are not overlapping
55   * Limitations
56     * If any object is at the boundary of the image, it will be rejected.
57     * This is because, if it's at the boundary, the midpoint calculated
58       is likely to be *very* inaccurate (because the object is likely to
59       have been cut off somewhat), so, instead of dealing with that
60       headache, it's just ignored.
61     * If the x midpoint of an object in both images is identical, that
62       object will be ignored because python doesn't like dividing by 0.
63   * For each of the objects identified in every frame
64     * Print their distance (in terms of Z depth) from the cameras, in metres.
65     * Along with their full X, Y, Z position from the cameras
66   * Assuming
67     * X baseline of 3500 metres
68     * Y baseline of 0
69     * Focal length of 12 metres
70     * Pixel spacing of 10 microns
71   * Limitations
72     * If an object is not present in both images, it shall be ignored, due
73       to not having full information about the x disparity for that frame.
74     * If an object is at the edge of one of the images, it will also be
75       ignored, due to the lack of accurate information about midpoints
76   * Work out the trajectories of each of the objects, printing a space-delimited
77     list of the identifiers of the objects that are probably UFOs (not moving
78     in a straight line)
79     * Assuming
80       * The object's position could be worked out for at least 3 frames.
81       * If fewer than 3 positions are known, there won't be enough points
82         in the line for the line to actually bend, so it'll treat it like
83         an asteroid.
84   * This is calculated via (ab)use of the dot products of unit vectors.
85     * Gets a normalized version of the vector between the first position and
86       the last position of the object in 3D space during these frames
87     * Also gets a normalized version of the vector that describes the
88       movement of the object between each 'frame'
89     * If this normalized vector is (0, 0, 0), we omit it because it will
90       completely mess up the maths.
91   * Finds the dot product of the normalized current frame movement vector

```

```

92         and the normalized total movement vector.
93         * Puts it on a list with all the found dot products
94     * Works out what 5% of the count of found dot products is
95     * Two unit vectors are equal if their dot product = 1.0
96         * Goes through that list of dot products, working out if that
97         dot product isClose to 1.0 (using isClose (hey!), to 9dp)
98         * We're using floats, and there's some inaccuracy with the
99         position measurement, so we know that we're not going to get
100         any dots that actually are going to be exactly 1.0
101     * If at least ~5% of the dots are **not** isClose to 1.0
102         * I am not ~95% sure that the line is straight, so, I'll accept the
103         null hypothesis that the object in question is a UFO.
104
105
106 ---
107
108 **USAGE**
109
110 ---
111
112 * python3 ce316ass.py 50 left-%03d.png right-%03d.png
113     * Assuming you have images called 'left-000.png' numbered to 'left-049.png',
114     and 'right-000.png' numbered to 'right-049.png' in this directory,
115     following all the assumptions/constraints in the 'PURPOSE' thing,
116     this will run.
117
118 ---
119
120 **AUTHOR**
121
122 --
123
124 Student 1804170
125
126 All the code written within this program is entirely my own work.
127
128 --
129
130 **RESULTS**
131
132 --
133
134 Given the sample data, this program produces an output of:
135     UFO: cyan white blue yellow orange
136
137 I was expecting Cyan to be a UFO. However, I wasn't really expecting the others
138 to be UFOs.
139
140 This is using a 'straightLineMaxUncertainty' global variable, which is defined
141 just above the 'isThisAStraightLine' function. Basically, if that proportion of
142 unit vector versions of the movement of the object between the frames are not
143 close enough to the unit vector of the start position to end position movement
144 (worked out 'if the dot of current movement and total movement = 1', because yay
145 dot product abuse), we can't confidently say it's a straight line, therefore,
146 we'll assume it's an asteroid.
147
148 straightLineMaxUncertainty must be between 0.0 and 1.0 (inclusive). If there
149 are very few current movements, the minimum 'sus' threshold will be 1.
150
151 Here's some outputs of the program with different 'straightLineMaxUncertainty'
152 values (specifically, the lowest values where I noticed a change in the number
153 of UFOs that were output).
154
155     * 0
156     * UFO: cyan red white blue yellow orange
157     * 0.05
158     * UFO: cyan white blue yellow orange
159     * 0.0625
160     * UFO: cyan blue yellow orange
161     * 0.075
162     * UFO: cyan blue orange
163     * 0.11
164     * UFO: cyan blue
165     * 0.125
166     * UFO: cyan
167     * 0.35
168     * UFO:
169
170 Green was never detected as a UFO.
171
172 Then there was the problem of what threshold to use.
173
174 So we need to consider the context of the problem.
175
176 The problem is 'which of these things are aliens trying to attack earth and nick
177 our PDMS'. Which, to me, sounds like the sort of program where false negatives
178 are more dangerous than false positives.
179
180 Therefore, to minimize the chance of false positives, and also to satisfy the
181 self-declared stats person inside me, I am going to stick with the 95% 'is not
182 UFO' confidence-y threshold thing.

```

```

183 So basically, when you see 'UFO' on the printout, read '>5% not an asteroid'.
184 Because if I'm not 95% confident of it being an asteroid (by having more than 5%
185 of the dot products of normalized movements * overall normalized movement not be
186 close to 1), I'm going to accuse it of being a UFO.
187
188
189 Yes, I know, probably wasting shots with the thing that shoots the objects.
190
191 However, seeing as not all of the objects are listed as 'UFOs' when using this
192 threshold, I'm confident that I'm not getting *too* many false positives, so I'm
193 considering it to be good enough.
194
195
196 ---
197 It runs pretty quickly though which is nice I guess.
198
199 6 seconds on the lab VMs!
200
201 """
202
203
204 # and time for some stuff to be imported.
205
206 import sys
207 # we need the command line arguments, the ability to quit, and the error stream
208 from math import sqrt, isclose # we need these for some of the maths.
209 from typing import Tuple, List, Dict, Union # No excuse to not type annotate.
210
211 import cv2
212 # we're doing computer vision, so opencv is also pretty darn useful.
213 import numpy as np # and it involves some numpy.ndarray objects!
214
215 """
216 --DEBUGGING FUNCTIONS FOR SEEING HOW THINGS WORK--
217
218 You ever wanted to see what goes on under the hood with this program?
219 No?
220
221 Either way, here are some functions that can be used for debugging.
222
223 Some code later on does have inbuilt debugging functions.
224 """
225
226 debugging: bool = False
227 """
228 Set this to 'True' if you want to enable these debug functions.
229 Or just keep it as 'False' if you want to just run the thing.
230 """
231
232
233 def debug(leftFilename: str, rightFilename: str,
234          leftIm: np.ndarray, rightIm: np.ndarray, frameNum: int = 0) -> None:
235     """
236     The wrapper function for the debugging functions, called by the main program
237     if 'debug' is set to True.
238
239     :param leftFilename: filename of the left image
240     :param rightFilename: filename of the right image
241     :param leftIm: the left image, BGR format
242     :param rightIm: the right image, BGR format
243     :param frameNum: What frame this is (used for labelling the mask
244     previews). If not defined, defaults to 0.
245     :return: nothing.
246     """
247
248     # prints the filenames, to make sure they're correct
249     print(leftFilename)
250     print(rightFilename)
251
252     # shows the left and right images
253     cv2.imshow("left", leftIm)
254     cv2.imshow("right", rightIm)
255     handleShowingStuff()
256     showMasksForDebugging(leftIm, rightIm, frameNum)
257     # and then proceeds to show all the masks produced for each object in
258     # the left and the right images.
259
260
261 # noinspection PyPep8Naming
262 def getObjectMasks(hsvIn: np.ndarray) -> Tuple[np.ndarray, np.ndarray,
263                                               np.ndarray, np.ndarray,
264                                               np.ndarray, np.ndarray,
265                                               np.ndarray]:
266     """
267     Generates the object masks for a given HSV image.
268
269     Yes, this uses some tuples that are defined as global variables in the
270     'actually important code' section after this section of debugging code,
271     because those tuples are part of the 'actually important code'.
272     This has been done so, if this debugging code section is removed (along with
273     the one branch in the main program that may potentially call this code),

```

```

274
275 Either way, this function won't be run until after those declarations are
276 reached, so no harm no foul or something along those lines.
277
278 This code is explained better in the 'getObjectMidpoints' function,
279 which has code that's basically identical to this.
280
281 You're probably going to ask 'why have you duplicated the code?'
282
283 Short answer: I don't want the production code to have a dependency on
284 this debug code, and, if this debug code was to be omitted, I'd have a
285 redundant dependency in the production code.
286
287 :param hsvIn: the hsv image that contains the objects
288 :return: masks for each of the 7 coloured objects that are in it.
289 (cyan, red, white, blue, green, yellow, orange)
290 """
291 cyan_mask: np.ndarray = cv2.inRange(hsvIn, min_cyan, max_cyan)
292 red_mask: np.ndarray = cv2.inRange(hsvIn, min_red, max_red)
293 white_mask: np.ndarray = cv2.inRange(hsvIn, min_white, max_white)
294 blue_mask: np.ndarray = cv2.inRange(hsvIn, min_blue, max_blue)
295 green_mask: np.ndarray = cv2.inRange(hsvIn, min_green, max_green)
296 yellow_mask: np.ndarray = cv2.inRange(hsvIn, min_yellow, max_yellow)
297 orange_mask: np.ndarray = cv2.inRange(hsvIn, min_orange, max_orange)
298
299 cyan_mask = cv2.morphologyEx(cyan_mask, cv2.MORPH_CLOSE, kernel33)
300
301 white_mask = cv2.subtract(white_mask, cyan_mask)
302
303 orange_mask = cv2.morphologyEx(orange_mask, cv2.MORPH_CLOSE, kernel33)
304
305 yellow_mask = cv2.subtract(yellow_mask, orange_mask)
306
307 return (cyan_mask, red_mask, white_mask, blue_mask,
308         green_mask, yellow_mask, orange_mask)
309
310 # noinspection PyPep8Naming
311 def getStereoMasks(left_in: np.ndarray, right_in: np.ndarray) -> \
312     List[Tuple[np.ndarray, np.ndarray]]:
313     """
314     Generates a list of masks for the left and right stereo images
315
316     :param left_in: the left image.
317     :param right_in: the right image.
318     :return: a list of tuples containing the masks for each of the images.
319     Tuples are in the form (left image mask, right image mask).
320     The list itself is in the order
321     cyan->red->white->blue->green->yellow->orange.
322     """
323     leftMasks: Tuple[np.ndarray, np.ndarray, np.ndarray, np.ndarray,
324                      np.ndarray, np.ndarray, np.ndarray] = \
325         getObjectMasks(cv2.cvtColor(left_in, cv2.COLOR_BGR2HSV))
326     rightMasks: Tuple[np.ndarray, np.ndarray, np.ndarray, np.ndarray,
327                      np.ndarray, np.ndarray, np.ndarray] = \
328         getObjectMasks(cv2.cvtColor(right_in, cv2.COLOR_BGR2HSV))
329
330     theMasks: List[Tuple[np.ndarray, np.ndarray]] = []
331     for i in range(0, 7):
332         theMasks.append((leftMasks[i], rightMasks[i]))
333
334     return theMasks
335
336
337 # noinspection PyPep8Naming
338 def showMasksForDebugging(i_left: np.ndarray, i_right: np.ndarray, f: int) -> \
339     None:
340     """
341     This is here mostly for debugging purposes, showing the masks of each image.
342     The 'left' image will have some text on it with the colour name identifier
343     of the object, as well as what number frame this is.
344     Apologies in advance if the annotations overlap with the position of
345     one of the objects in the set of images you are using to mark this.
346
347     :param i_left: left image (BGR format)
348     :param i_right: right image (BGR format)
349     :param f: frame number
350     :return: nothing.
351     """
352     # order in which the masks of objects are returned
353     objectOrder: List[str] = \
354         ["cyan", "red", "white", "blue", "green", "yellow", "orange"]
355
356     i: int = 0
357     debugMasks: List[Tuple[np.ndarray, np.ndarray]] = \
358         getStereoMasks(i_left, i_right) # the masks for each image.
359     for m in debugMasks:
360         showLeft: np.ndarray = np.ndarray.copy(m[0]) # copy of the left mask
361         label: str = str(objectOrder[i]) + " " + str(f)
362         print(label)
363         # puts some text with object id and frame num on the left mask copy.
364

```

```

365     cv2.putText(showLeft, label, (50, 50),
366                 cv2.FONT_HERSHEY_SIMPLEX, 1, 255, 2, cv2.LINE_AA)
367     # and shows the left/right masks
368     cv2.imshow("left", showLeft)
369     cv2.imshow("right", m[1])
370     handleShowingStuff()
371
372     # finds midpoints of the object on the mask
373     leftMid: Tuple[float, float] = getObjectMidpoint(m[0])
374     rightMid: Tuple[float, float] = getObjectMidpoint(m[1])
375
376     print(leftMid)
377     print(rightMid)
378
379     showLeft: np.ndarray = cv2.bitwise_and(i_left, i_left, mask=m[0])
380     showRight: np.ndarray = cv2.bitwise_and(i_right, i_right, mask=m[1])
381     # puts the annotation on the left copy again.
382     cv2.putText(showLeft, label, (50, 50),
383                 cv2.FONT_HERSHEY_SIMPLEX, 1, (255, 255, 255), 2,
384                 cv2.LINE_AA)
385     if leftMid != (-1, -1):
386         # if the left midpoint is valid, it's put on the left copy as a
387         # magenta dot at the int version of the midpoint.
388         iLeftMid: Tuple[int, int] = (int(leftMid[0]), int(leftMid[1]))
389         cv2.line(showLeft, iLeftMid, iLeftMid, (255, 0, 255), 1)
390     if rightMid != (-1, -1):
391         # ditto for the right image.
392         iRightMid: Tuple[int, int] = (int(rightMid[0]), int(rightMid[1]))
393         cv2.line(showRight, iRightMid, iRightMid, (255, 0, 255), 1)
394
395     # and shows them.
396     cv2.imshow("left", showLeft)
397     cv2.imshow("right", showRight)
398     handleShowingStuff()
399     i += 1
400
401
402 # noinspection PyPep8Naming
403 def handleShowingStuff() -> None:
404     """
405     This is here to handle actually showing stuff when debugging the program.
406     Prints something to console to let the user know that they need to press
407     something to continue, handles doing the cv2.waitKey(0) call (thereby
408     allowing any pending cv2.imshow()'d images to be shown), and,
409     if q or escape are pressed, the program will close.
410
411     :return: nothing is returned.
412     """
413     print("press something to continue (press q or escape to quit)")
414     key: int = cv2.waitKey(0)
415     # quit if escape (27) or q (113) are pressed
416     if key == 27 or key == 113:
417         cv2.destroyAllWindows()
418         print("quitting!")
419         sys.exit(0)
420
421
422 """
423 -- THE ACTUALLY IMPORTANT CODE THAT ACTUALLY DOES STUFF --
424
425 Yep. Everything from here is actually of some use.
426 """
427
428 """
429 ~~~~~ A VECTOR3D CLASS (and also a function that uses it) ~~~~~
430
431 This is used later on, represents points in 3D space
432 """
433
434
435 class Vector3D:
436     """
437     A class to represent a vector in 3D space.
438     This code was written by myself, but I fully acknowledge that somebody else
439     has probably written a python implementation of a 3D vector before, so any
440     relationship between this Vector3D and another implementation of Vector3D
441     is entirely coincidental.
442
443     This implementation just contains a normalize, subtract, isZero, dot product
444     , and __str__ method (as well as a constructor ofc), because that's all
445     the math stuff I needed for this particular use case.
446
447     References for particular sources for math stuff (when used) have been given
448     in the methods for each of the functions that use the math stuff.
449
450     Now, you may ask yourself 'why is a class being used, when a tuple could
451     do the same thing?'. Simple answer; encapsulation (so I have the methods all
452     in the same place as each other). And also making sure I don't get confused
453     between tuples of floats and 3D vectors. The mutability is also nice for the
454     subtraction and the normalization stuff.
455     """

```

```

456 def __init__(self, x: float, y: float, z: float):
457     """
458     Constructs a Vector3D with the given x, y, and z coordinates
459     :param x: x coordinate
460     :param y: y coordinate
461     :param z: z coordinate
462     """
463     self.x: float = x
464     self.y: float = y
465     self.z: float = z
466
467 def magnitude(self) -> float:
468     """
469     dist between 2 3D points:
470     sqrt(((x2-x1)^2) + ((y2-y1)^2) + ((z2 - z1)^2))
471     and we know that x1,y1,z1 = 0 already (because vector comes from origin
472     0) and x2, y2, and z2 are the x, y, and z of this vector.
473
474     Got the maths from https://www.calculator.net/distance-calculator.html
475
476     :return: the magnitude of this vector
477     """
478     return sqrt((self.x ** 2) + (self.y ** 2) + (self.z ** 2))
479
480 def normalized(self) -> "Vector3D":
481     """
482     Normalizes this Vector3D (makes the magnitude 1 by dividing all the
483     components of this Vector3D by its magnitude)
484
485     :return: This Vector3D, but with a magnitude of 1 instead.
486     if this already had a magnitude of 0, it'll return itself as-is.
487     """
488     mag: float = self.magnitude()
489     if mag > 0:
490         self.x = self.x / mag
491         self.y = self.y / mag
492         self.z = self.z / mag
493
494     return self
495
496 def subtract(self, other: "Vector3D") -> "Vector3D":
497     """
498     Subtracts the other Vector3D from this Vector3D, returning this
499     modified Vector3D.
500
501     Didn't need to get the maths from anywhere because subtraction is pretty
502     darn simple and doesn't have any weirdness.
503
504     :param other: the other Vector3D to subtract from this.
505     :return: this Vector3D minus 'other'. Would have type-annotated the
506     return type as Vector3D, but python didn't like that.
507     """
508     self.x -= other.x
509     self.y -= other.y
510     self.z -= other.z
511     return self
512
513 def dot(self, other: "Vector3D") -> float:
514     """
515     Returns the dot product of this Vector3D and the other Vector3D.
516
517     Got the maths from https://www.quantumstudy.com/physics/vectors-2/
518
519     :param other: the other Vector3D this is being dot product-ed against.
520     :return: the dot product of this and the other vector3D
521     """
522     return (self.x * other.x) + (self.y * other.y) + (self.z * other.z)
523
524 # noinspection PyPep8Naming
525 def isZero(self) -> bool:
526     """
527     Check if this vector is (0,0,0)
528     :return: Returns true if x, y and z are exactly equal to 0
529     """
530     return (self.x == 0) and (self.y == 0) and (self.z == 0)
531
532 def __str__(self) -> str:
533     """
534     Outputs this as a string; as a tuple in the form (x, y, z)
535     :return: a string of the tuple (self.x, self.y, self.z)
536
537     """
538     return str((self.x, self.y, self.z))
539
540
541 # noinspection PyPep8Naming
542 def normalizeVectorBetweenPoints(fromVec: Vector3D, toVec: Vector3D) ->\
543     Vector3D:
544     """
545     Get lhs-rhs but normalized instead (leaving lhs and rhs untouched)
546

```

```

547 >>> normalizeVectorBetweenPoints(Vector3D(0,0,0),Vector3D(1,1,1))
548 (0.5773502691896258, 0.5773502691896258, 0.5773502691896258)
549
550 >>> normalizeVectorBetweenPoints(Vector3D(0,0,0),Vector3D(2,2,2))
551 (0.5773502691896258, 0.5773502691896258, 0.5773502691896258)
552
553 >>> normalizeVectorBetweenPoints(Vector3D(0,0,0),Vector3D(1,1.5,2))
554 (0.3713906763541037, 0.5570860145311556, 0.7427813527082074)
555
556 >>> normalizeVectorBetweenPoints(Vector3D(1,1,1),Vector3D(1,1,1))
557 (0, 0, 0)
558
559 :param fromVec: going from this vector
560 :param toVec: to this other vector
561 :return: toVec - fromVec but normalized. Or in other words, the direction of
562 movement from the position 'fromVec' to the position 'toVec'
563 """
564 return Vector3D(toVec.x, toVec.y, toVec.z)\
565     .subtract(fromVec)\
566     .normalized()
567
568
569 """
570 ~~~~~ READING IMAGES, FINDING OBJECTS, AND ALSO CALCULATING AND
571 PRINTING THE POSITIONS OF SAID OBJECTS ~~~~~
572
573 These functions (and also globals) are responsible for finding and printing the
574 positions of the objects in 3D space.
575 """
576
577
578 # HSV values to be used to extract the objects from the image.
579 min_cyan: Tuple[int, int, int] = (90, 0, 0)
580 """minimum HSV threshold for the cyan object"""
581 max_cyan: Tuple[int, int, int] = (90, 255, 255)
582 """maximum HSV threshold for the cyan object"""
583
584 min_red: Tuple[int, int, int] = (0, 1, 0)
585 """minimum HSV threshold for the red object"""
586 max_red: Tuple[int, int, int] = (0, 255, 255)
587 """maximum HSV threshold for the red object"""
588
589 min_white: Tuple[int, int, int] = (0, 0, 1)
590 """minimum HSV threshold for the white object"""
591 max_white: Tuple[int, int, int] = (0, 0, 255)
592 """maximum HSV threshold for the white object"""
593
594 min_blue: Tuple[int, int, int] = (120, 1, 1)
595 """minimum HSV threshold for the blue object"""
596 max_blue: Tuple[int, int, int] = (120, 255, 255)
597 """maximum HSV threshold for the blue object"""
598
599 min_green: Tuple[int, int, int] = (60, 1, 1)
600 """minimum HSV threshold for the green object"""
601 max_green: Tuple[int, int, int] = (60, 255, 255)
602 """maximum HSV threshold for the green object"""
603
604 min_yellow: Tuple[int, int, int] = (30, 1, 0)
605 """minimum HSV threshold for the yellow object"""
606 max_yellow: Tuple[int, int, int] = (30, 255, 255)
607 """maximum HSV threshold for the yellow object"""
608
609 min_orange: Tuple[int, int, int] = (20, 0, 0)
610 """minimum HSV threshold for the orange object"""
611 max_orange: Tuple[int, int, int] = (29, 255, 255)
612 """maximum HSV threshold for the orange object"""
613
614 kernel33: np.ndarray = np.ones((3, 3), np.uint8)
615 """
616 A 3*3 numpy array of 1s, to be used when closing up holes in some of the masks
617 """
618
619
620 # noinspection PyPep8Naming
621 def getObjectMidpoint(objectMask: np.ndarray) -> Tuple[float, float]:
622     """
623     Returns the midpoint of the region of 1s in the given binary image array
624
625     If there is no region of 1s, we return (-1,-1). If there's a 1 on the
626     boundary of the image, once again, we return (-1,-1), because we'll know
627     that the midpoint we find will probably be inaccurate.
628
629     This estimates the midpoint by finding the upper/lower X and Y bounds of
630     that region of 1s in the image. Yes, it's a pretty naive, brute force-y
631     method. However, I tried several object/blob detection algorithms within
632     OpenCV, however, none of them really worked as intended (not detecting the
633     objects in the earlier images, not really being able to work out which
634     keypoints correspond to each object, and refusing to detect the single
635     object when I go through the effort of masking out everything else in the
636     image), so I went 'fuck it, guess I'm doing it myself'
637

```



```

638 DISCLAIMER: I produced this code on the 3rd of March, several days before
639 that email was sent out suggesting that a procedure like this would be
640 worth using for the assignment.
641
642 :param objectMask: binary image, with 1s in the area where the object
643 with the midpoint being looked for is, and 0s everywhere else.
644 :return: A tuple holding the the midpoint of the object.
645 If the object isn't present (mask all 0s), a value of (-1,-1) is returned.
646 Additionally, if the object is at the edge of the image (a minimum is 0,
647 or a maximum is at the maximum possible x/y), that heavily implies that
648 the object is partially out-of-frame. Therefore, as that means the true
649 bounds are likely to be out-of-frame, this midpoint detector will not find
650 the true midpoint of the object, so it will give up and return -1s for that
651 as well.
652 """
653
654 if objectMask.any():
655     if objectMask[0].any() or objectMask[-1].any():
656         # if there's anything in the topmost or bottommost row, that means
657         # there's something on the image boundary, meaning that the midpoint
658         # found will be inaccurate, so we're not going to bother finding it.
659         return -1, -1
660     else:
661         # if nothing in the objectMask is a 1, we return -1s.
662         return -1, -1
663
664 # obtaining the shape of the actual mask image
665 yx: Tuple[int, int] = objectMask.shape
666 ny: int = yx[0]
667 nx: int = yx[1]
668
669 # and declaring some variables to hold the info we find out about
670 # the shape of the object.
671 minX: int = -1
672 maxX: int = -1
673 minY: int = -1
674 maxY: int = -1
675 notFoundFirst: bool = True # set this to false when we find first pixel
676
677 # now we just casually loop through the image pixels,
678 # and find out about what sort of shape the object has
679 for y in range(1, ny-1):
680     # we already established that the topmost/bottommost rows are empty.
681     if objectMask[y].any():
682         # we only bother with this row if it contains 1s
683         maxY = y
684         # y wont get smaller.
685         # and assignment has same complexity as checking a single
686         # condition so I may as well just reassign y anyway.
687         for x in range(0, nx):
688             if objectMask[y][x] != 0:
689                 # if it's not 0, we've found something!
690                 if notFoundFirst:
691                     # if we haven't found the first thing yet, we have now.
692                     notFoundFirst = False
693                     minX = maxX = x
694                     minY = maxY = y
695                 else:
696                     if x < minX:
697                         minX = x
698                     elif x > maxX:
699                         maxX = x
700
701 # if it's at the x bounds of the image, the result definitely won't be
702 # accurate, so we'll just return -1, -1.
703 # (we already checked the y bounds earlier on)
704 if (minX == 0) or (maxX == nx-1):
705     return -1, -1
706
707 # working out widths and heights
708 w: int = (maxX - minX)
709 h: int = (maxY - minY)
710
711 # using that and the lower bounds for x and y to find the midpoints
712 xMid: float = minX + (w / 2.0)
713 yMid: float = minY + (h / 2.0)
714
715 # and returning a tuple with those midpoints
716 return xMid, yMid
717
718
719 # noinspection PyPep8Naming
720 def getObjectMidpoints(hsvIn: np.ndarray) -> Dict[str, Tuple[float, float]]:
721     """
722     Gets the midpoints for the objects that may be in the given HSV image.
723     We generate masks that contain only the region of the HSV object that is
724     occupied by the pixels that make up a particular object, do a bit of cleanup
725     for the objects that have somewhat overlapping pixel values, and then
726     put those masks into getObjectMidpoint to produce a dictionary holding the
727     midpoints of each of the objects in the image.
728 """

```

```

729 If an object is not present in the image, its entry in the dictionary will
730 have the default value of (-1,-1) instead of an actual midpoint.
731
732 You might be wondering 'why am I making all the masks at once and then
733 finding the midpoints from them all at once instead of just making a mask,
734 getting the midpoint, and moving on to the next mask?'
735
736 Simple answer: Doing it this way allows me to do the cleanup stuff that
737 needs to be done for cyan/white/orange/yellow more effectively,
738 and it means I can declare the results dictionary as a literal and also
739 immediately return it. yay efficiency.
740
741 :param hsvIn: the hsv image that contains the objects
742 :return: dict with midpoints for each of the 7 coloured objects
743         that might be present in the given image. If not present, midpoint will
744         be (-1,-1). Keys are (cyan, red, white, blue, green, yellow, orange)
745 """
746 # generating masks for each object.
747 cyan_mask: np.ndarray = cv2.inRange(hsvIn, min_cyan, max_cyan)
748 red_mask: np.ndarray = cv2.inRange(hsvIn, min_red, max_red)
749 white_mask: np.ndarray = cv2.inRange(hsvIn, min_white, max_white)
750 blue_mask: np.ndarray = cv2.inRange(hsvIn, min_blue, max_blue)
751 green_mask: np.ndarray = cv2.inRange(hsvIn, min_green, max_green)
752 yellow_mask: np.ndarray = cv2.inRange(hsvIn, min_yellow, max_yellow)
753 orange_mask: np.ndarray = cv2.inRange(hsvIn, min_orange, max_orange)
754
755 # and now, time for some cleanup
756
757 # filling in the midpoint for the cyan mask (as that's actually white)
758 cyan_mask = cv2.morphologyEx(cyan_mask, cv2.MORPH_CLOSE, kernel33)
759
760 # removing the midpoint in the cyan mask from the white mask
761 white_mask = cv2.subtract(white_mask, cyan_mask)
762
763 # filling in the midpoint for the orange mask (as that's actually yellow)
764 orange_mask = cv2.morphologyEx(orange_mask, cv2.MORPH_CLOSE, kernel33)
765
766 # and removing the orange midpoint from some yellow
767 yellow_mask = cv2.subtract(yellow_mask, orange_mask)
768
769 # finally making + returning the dict with the midpoints of each object
770 return {
771     "cyan": getObjectMidpoint(cyan_mask),
772     "red": getObjectMidpoint(red_mask),
773     "white": getObjectMidpoint(white_mask),
774     "blue": getObjectMidpoint(blue_mask),
775     "green": getObjectMidpoint(green_mask),
776     "yellow": getObjectMidpoint(yellow_mask),
777     "orange": getObjectMidpoint(orange_mask)
778 }
779
780 # noinspection PyPep8Naming
781 def getStereoPositions(left_in: np.ndarray, right_in: np.ndarray) -> \
782     Dict[str,
783         Tuple[Tuple[float, float],
784             Tuple[float, float]]]:
785     """
786     Gets the positions of objects in the left and right *coloured* images.
787     This **will** assume that the dimensions of the left and right images
788     are identical, and that the two images are Y-aligned already.
789
790     :param left_in: the left input image *in colour*
791     :param right_in: the right input image *in colour*
792     :return: a dictionary with the names of the objects in both images as keys,
793             and a tuple, containing the x' and y' positions of that particular object
794             in both images as the value
795
796             1st tuple in the value tuple: (x',y') from left image.
797             2nd tuple in the value tuple: (x',y') from right image.
798
799             If an object is not present in **both** images, it will **not** be
800             present in the returned dictionary.
801     :raises: ValueError if the two images provided have different dimensions.
802     """
803     # gets the shape of the images
804     yx: Tuple[int, int] = left_in.shape
805     if yx != right_in.shape:
806         # complains if the dimensions aren't identical.
807         raise ValueError("Please provide images with identical dimensions.")
808
809     lDict: Dict[str, Tuple[float, float]] = getObjectMidpoints(left_in)
810     """dictionary with midpoints for every object in the left image"""
811
812     rDict: Dict[str, Tuple[float, float]] = getObjectMidpoints(right_in)
813     """dictionary with midpoints for every object in the right image"""
814
815     # half of the x and y dimensions of the images
816     halfY: float = yx[0] / 2
817     halfX: float = yx[1] / 2
818
819

```

```

820 posDict: Dict[str, Tuple[Tuple[float, float], Tuple[float, float]]] = {}
821 """a dictionary for all the calculated (X',Y') positions for each image"""
822
823 # obtains the keys from lDict but as a list so it can be foreach'd,
824 # and also foreaches through them
825 for key in [*lDict.keys()]:
826     if lDict[key] != (-1, -1):
827         if rDict[key] != (-1, -1):
828             # if both dictionaries have an actual value for the key
829
830             # just getting a copy of those raw values real quick
831             rawL: Tuple[float, float] = lDict[key]
832             rawR: Tuple[float, float] = rDict[key]
833
834             # work out the X' and the Y' stuff for left and right
835             # and put it into the posDict.
836             # X' = x - halfX
837             # Y' = halfY - y
838             posDict[key] = (
839                 (rawL[0] - halfX, halfY - rawL[1]),
840                 (rawR[0] - halfX, halfY - rawR[1])
841             )
842
843 # and now return the posDict
844 return posDict
845
846
847 placeholderOutString: str = "{:5} {:8} {:.2e} {}"
848 """
849 This is a placeholder string to be used when formatting the frame-by-frame
850 printout of object data. Double space between each thing of data.
851 1st value: will be the frame number. 5 width, right-aligned.
852 2nd value: object identifier. 8 width, also left-aligned
853 3rd value: object distance (Z pos, in metres). 8 width, in the form 1.23e+45
854
855 4th value: just the raw (X,Y,Z) position of the object in 3D space (in metres),
856 for sake of curiosity.
857 """
858 focalLength: float = 12
859 "Focal length of camera is 12m"
860 baseline: float = 3500
861 "baseline between cameras is 3.5km -> 3500m"
862 pixelSize: float = float(1e-5)
863 "pixel spacing: 10 microns -> 1e-5 metres"
864
865
866 # noinspection PyPep8Naming
867 def calculateAndPrintPositionsOfObjects(leftIm: np.ndarray,
868                                         rightIm: np.ndarray,
869                                         frameNum: int = 0) -> \
870     Dict[str, Vector3D]:
871     """
872     Given a left image (BGR), a right image (BGR), and a frame number (optional)
873     , this method will print the details about the identifiers and the depths
874     (Z axis positions) of the objects in the image (formatted as per the
875     assignment brief, using the global placeholderOutString), and will return a
876     dictionary with vector3 positions of the objects in the images.
877
878     Objects in only one image will be omitted. Distances will be in metres.
879
880     leftIm and rightIm must be in BGR colour, and have identical dimensions.
881
882     This will use the focalLength, baseline, and pixelSize global variables to
883     calculate the positions of the objects.
884
885     placeholderOutString, focalLength, baseline, and pixelSize are present
886     just above this function.
887
888     :param leftIm: The left stereo image to look at (BGR colour)
889     :param rightIm: The right stereo image to look at (BGR colour)
890     :param frameNum: The frame number (optional). Will only be used to prefix
891     the printout. If not supplied, 0 will be used.
892     :return: A dictionary with the identifiers of the objects identified, along
893     with their vector3 positions, in metres, relative to the midpoint between
894     the cameras.
895     """
896
897     imgPositions: \
898         Dict[str, Tuple[Tuple[float, float], Tuple[float, float]]] = \
899         getStereoPositions(cv2.cvtColor(leftIm, cv2.COLOR_BGR2HSV),
900                             cv2.cvtColor(rightIm, cv2.COLOR_BGR2HSV))
901     """
902     These are the left and right image X'Y' coords of every object
903     in both the left and the right images.
904     """
905
906     # noinspection PyShadowingNames
907     posXYZ: Dict[str, Vector3D] = {}
908     """
909     this will hold the X, Y, Z coords of the objects in 3D space,
910     using the midpoint between the cameras as the origin,

```

```

911     measured in metres.
912     """
913
914     # unpacking the keys/object names as a list so we can iterate through them.
915     # Why do I need to do this? Because Dict.keys() returns a KeysView object,
916     # which isn't iterable, and is generally awkward to work with. However,
917     # putting a KeysView kv into [*kv] basically unpacks it into a list, which
918     # we can iterate through. So that's what happens here.
919     for key in [*imgPositions.keys()]:
920
921         currentPos: Tuple[Tuple[float, float], Tuple[float, float]] = \
922             imgPositions[key]
923         "We obtain info about current object's 2d pos from imagePositions"
924
925         xDisparity: float = currentPos[0][0] - currentPos[1][0]
926         "x disparity = xL - xR"
927
928         if xDisparity == 0.0:
929             # undefined behaviour (aka runtime error) if the x disparity is 0.
930             # so we'll just forget it ever happened
931             continue
932
933         rawZ: float = (focalLength * baseline) / (xDisparity * pixelSize)
934         """
935         Z = (f * b) / (xL - xR)
936         which is how we work out what dist is.
937         """
938
939         rawX: float = ((-currentPos[1][0] * pixelSize) / focalLength) * rawZ
940         """
941         (-xr/f) = (X/Z), therefore (-xr/f) * Z = X
942         So I used that equation to find out what the actual X position
943         of the object is in 3D space.
944         """
945
946         yMid: float = (currentPos[0][1] + currentPos[1][1]) / 2
947         """
948         This is the y midpoint of the object. I'm getting the average of the y
949         position for the two images, just in case they differ a bit (and, if
950         they're actually identical, the yMid will just be the same as them)
951         """
952
953         rawY: float = -((yMid * pixelSize) / focalLength) * rawZ
954         """
955         (xL / f) = (B-X)/Z
956         so, substituting the x for y
957         (y / f) = (B-Y)/Z = -Y/Z
958         and if we rearrange that so -Y is the result
959         (y/f) * Z = -Y
960         and negating that to get positive Y as the result
961         -(y/f) * Z = Y
962         """
963
964         # we put the raw XYZ into posXYZ
965         posXYZ[key] = Vector3D(rawX, rawY, rawZ)
966
967         # Now, we just print the required info, as per the specification.
968         print(placeholderOutString.format(
969             frameNum, # what frame number this is
970             key, # identifier of this object
971             rawZ, # we print the Z depth
972             posXYZ[key] # the XYZ pos, printed for debug reasons.
973         ))
974
975     # we finish by returning the posXYZ dictionary.
976     return posXYZ
977
978     """
979     ~~~~~~ WORKING OUT WHAT IS/IS NOT A UFO ~~~~~~
980
981     These methods (and globals) are used to work out what is/isn't a UFO from a
982     dictionary of UFO identifiers and their frame-by-frame positions as lists of
983     Vector3D objects.
984     """
985
986
987     debuggingLineStuff: bool = False
988     """
989     Set this to true if you want to enable the debug printouts for the
990     isThisAStraightLine function (immediately below this)
991     """
992
993     straightLineMaxUncertainty: float = 0.05
994     """
995     How uncertain we are allowing ourselves to be about whether a line is straight
996     or not. If there is more than this amount of uncertainty (0.05 = 5%), we won't
997     consider it to be a straight line; instead, we'll consider it to be a UFO.
998
999     MUST BE BETWEEN 0 AND 1.0!
1000
1001     Outputs at different thresholds of this value:

```

```

1002     * 0
1003     * UFO: cyan red white blue yellow orange
1004     * 0.05
1005     * UFO: cyan white blue yellow orange
1006     * 0.0625
1007     * UFO: cyan blue yellow orange
1008     * 0.075
1009     * UFO: cyan blue orange
1010     * 0.11
1011     * UFO: cyan blue
1012     * 0.125
1013     * UFO: cyan
1014     * 0.35
1015     * UFO:
1016
1017 """
1018
1019 assert (0.0 <= straightLineMaxUncertainty <= 1.0)
1020
1021
1022 # noinspection PyPep8Naming
1023 def isThisAStraightLine(line: List[Vector3D]) -> bool:
1024     """
1025     Returns whether or not a sequence of 3D points is a straight line,
1026     using the getNormDifferenceBetweenPoints function, and dot product abuse.
1027
1028     If there's 2 or fewer points, it certainly ain't bent, so it will return
1029     true.
1030
1031     Due to the inherent uncertainty with how the points are calculated, I am
1032     giving some leeway in the calculations.
1033
1034     And basically I'm working out if it's straight or not by seeing if at least
1035     ~95% of the dot products for the normalized vectors between each vector of
1036     the line and the normalized vector between the start and the end of the line
1037     are ~1.0 (tl;dr the dot product of two identical unit vectors is 1, but, if
1038     they aren't identical, it'll be less than 1).
1039
1040     :param line: the sequence of 3D points
1041     :return: true if they're a straight enough line, false otherwise.
1042     """
1043
1044     if len(line) < 3:
1045         # 3 short 5 bend
1046         return True
1047
1048     startEndDiff: Vector3D = \
1049         normalizeVectorBetweenPoints(line[0], line[-1])
1050     """
1051     This is a normalized vector between the starting point and the ending
1052     point of the object. Every single vector between each pair of consecutive
1053     points will be checked for similarity to this via their dot products
1054     """
1055
1056     if debuggingLineStuff:
1057         print(startEndDiff)
1058
1059     dots: List[float] = []
1060     """
1061     A list to hold the dot product(s) of startEndDiff and the normalized
1062     versions of the vectors between each pair of consecutive vectors in line
1063     """
1064
1065     thisIndex: int = 0
1066     """
1067     A cursor to the index of the line used for this iteration. This starts at 0,
1068     so, when the first iteration increments it to 1, the first iteration will
1069     look at indexes [0] and [1] (getting the first movement vector).
1070     """
1071
1072     while True:
1073         thisIndex += 1 # we move to the next index of the list
1074         if thisIndex >= len(line):
1075             # we're basically emulating a do/while loop here
1076             # with a while condition of thisIndex < len(line)
1077             # so when we get to the end of the list, we stop looping.
1078             break
1079
1080         thisDiff: Vector3D = \
1081             normalizeVectorBetweenPoints(line[thisIndex-1], line[thisIndex])
1082         """
1083         We find the vector between the position at the index thisIndex and the
1084         point on the line behind it, but normalized instead, so we can compare
1085         it to the normalized startEndDiff.
1086         """
1087
1088         if debuggingLineStuff:
1089             print(thisDiff)
1090
1091         if thisDiff.isZero():
1092             # if it's a 0 vector, that will mess up our calculations, so

```

```

1093         # we'll just ignore it and move on to the next pair of vectors.
1094         continue
1095
1096     thisDot: float = startEndDiff.dot(thisDiff)
1097     """
1098     TIME FOR SOME ILLEGAL MATHS!!!
1099
1100     Funnily enough, you can actually use the dot product of two unit vectors
1101     to compare the unit vectors for similarity.
1102
1103     Unit vectors have a magnitude of 1. And the dot product of two vectors
1104     is basically working out how far a vector projects onto another. Forgot
1105     the technical terms.
1106
1107     But, the important thing is that if you have two unit vectors, and the
1108     two unit vectors are identical (same x, y, z; same direction), the dot
1109     product of those two vectors will be 1.
1110
1111     For a more practical example of this illegal maths in action,
1112     there's a rather nice demo of the dot products of vectors (but in two
1113     dimensions) here, where you can try messing around with unit vectors:
1114     https://www.youphysics.education/scalar-and-vector-quantities/dot-product/
1115     """
1116
1117     if debuggingLineStuff:
1118         print(thisDot)
1119
1120     dots.append(thisDot) # and we append the current dot product to dots.
1121
1122 if len(dots) == 0:
1123     # if all the differences between positions were (0,0,0), this ain't bent
1124     # so it'll return True.
1125     return True
1126
1127 maxSusFloat: float = len(dots) * straightLineMaxUncertainty
1128
1129 maxSus: int = int(maxSusFloat)
1130 """
1131 This is how many of the dot products have to be not roughly equal to 1 for
1132 the object to be labelled as a UFO. It's currently set up so, if ~5% of the
1133 dots are not equal to 1, that's sus enough for us to label it as a UFO, with
1134 95% certainty of this being the case.
1135
1136 This is because I'm working on the hypothesis that 'This line is straight',
1137 and I'm going only going to accept this hypothesis with a certainty of at
1138 least 95% (it's good enough for geography, and there's not enough data, at
1139 least in the sample dataset, for me to really be able to test for 99%
1140 certainty)
1141
1142 So, if ~5% of the vectors indicate that this is not travelling in a straight
1143 line, we have a certainty of less than 95% that this is travelling in a
1144 straight line, therefore, we will reject the hypothesis that 'this object
1145 is travelling in a straight line', and accept the null hypothesis (of 'this
1146 object is not travelling in a straight line') instead.
1147 """
1148
1149 # and if there's a maximum sus level that's below 0, we set it to 1.
1150 if maxSus < 1:
1151     maxSus = 1
1152
1153 if debuggingLineStuff:
1154     print("ufo is " + str(maxSus) + " sus!") # WHEN THE UFO IS SUS!
1155
1156 susCount: int = 0
1157 "The count of how many times the dot product has been not equal to 1."
1158
1159 debugCount: int = 0
1160 "just here as a printout for debug purposes."
1161
1162 for d in dots:
1163     if not isclose(1, d, rel_tol=1e-09):
1164         """
1165         We're using floating-point numbers here, so, because it's nigh
1166         impossible to get a dot of 1.0 (mostly due to the slight inaccuracy
1167         inherent due to resolution and pixel values and stuff like that),
1168         we're using the 'isclose' method to check if the dots are within
1169         1e-09 of 1 (at least 0.999999999).
1170
1171         If it isn't close enough to 1, the thing isn't going in
1172         a straight line. So it's sus.
1173
1174         And if it's sus maxSus times, this is clearly not a straight line.
1175         """
1176         susCount += 1
1177         if debuggingLineStuff:
1178             print("diff " + str(debugCount) + " is " +
1179                   str(susCount) + " sus")
1180         if susCount >= maxSus: # WHEN THE UFO IS SUS
1181             return False # amogus
1182         debugCount += 1

```

```

1184     # if it hasn't been thrown out as sus yet, it's probably a straight line.
1185     return True
1186
1187
1188 # noinspection PyPep8Naming
1189 def makeUfoString(objPositions: Dict[str, List[Vector3D]]) -> str:
1190     """
1191     Given the dictionary of object identifiers + lists with all of
1192     their Vector3D positions, create the space delimited string with the
1193     list of all the identifiers of objects that are UFOs
1194     :param objPositions: dictionary of object identifiers + Vector3D
1195     positions for all the objects that may or may not be UFOs
1196     :return: string with the identifiers of what is and isn't a UFO
1197     """
1198     ufoString: str = ""
1199     """
1200     The space-delimited string of UFO identifiers.
1201     """
1202
1203     for key in [*objPositions.keys()]:
1204         if not isThisAStraightLine(objPositions[key]):
1205             # we check if the list of points is a straight line.
1206             # If they aren't a straight line, we know this is a UFO, so
1207             # it's appended to the ufoString.
1208             ufoString = ufoString + " " + key
1209
1210     return ufoString
1211
1212
1213 """
1214 ~~~~~ CHECKING WHETHER OR NOT AN IMAGE OPENED ~~~~~
1215
1216 yeah this is just here to stop a big error from happening
1217 """
1218
1219
1220 # noinspection PyPep8Naming
1221 def checkIfImageWasOpened(filename: str, img: Union[np.ndarray, None]) -> None:
1222     """
1223     This will check if the image with the given filename could be opened
1224
1225     :param filename: the name of the file
1226     :param img: the numpy.ndarray (or lack thereof) that opencv could open
1227     using that given filename
1228     :return: nothing. But, if the file couldn't be opened (causing img to be
1229     None instead of a np.ndarray), the program complains and promptly closes.
1230     """
1231     if isinstance(img, type(None)):
1232         """
1233         If the image is actually NoneType, the program complains and closes.
1234
1235         And you may ask yourself 'why am I doing this in such a convoluted way?'
1236         Simple answer: There is no simpler way to do it.
1237         Opencv doesn't throw an exception if the image couldn't be read, it just
1238         returns None.
1239         So, if it does return None, I could just detect it with 'if im == None',
1240         right? WRONG!
1241         Thing is, if it doesn't return None, it returns a numpy.ndarray. And if
1242         you attempt to compare one of those against None, guess what? You get
1243         a ValueError and a snarky comment saying 'oh no the truth value of this
1244         is ambiguous pls use .any() or .all() instead'
1245         But if I try to use those methods to check if the image exists, and the
1246         image doesn't exist, guess what? You get an AttributeError.
1247
1248         Now, do I want to bother with throwing and catching exceptions manually?
1249         No. cba to deal with that overhead.
1250
1251         Would I have preferred it if OpenCV could have just thrown an exception
1252         or just returned an empty array instead of returning a mcfucking None?
1253
1254         Yes.
1255
1256         But, alas, we live in a society. Rant over.
1257         """
1258         print("ERROR: Could not open the file called " + filename)
1259         sys.exit(1)
1260
1261
1262 """
1263 ~~~~~ THE MAIN PROGRAM ~~~~~
1264
1265 Everything from here is the stuff that runs when you start running this.
1266 """
1267
1268
1269 if len(sys.argv) < 4:
1270     # If you don't give 3 command line arguments, the program will complain
1271     print("Usage:", sys.argv[0],
1272           "<frame count> ",
1273           "<left frame filename template, such as left-%03d.png> ",
1274           "<right frame filename template, such as right-%03d.png>",

```

```

1275         file=sys.stderr)
1276     # and promptly quit
1277     sys.exit(1)
1278
1279
1280 # this line was adapted from the assignment brief.
1281 nframes: int = int(sys.argv[1])
1282 """
1283 Reads the 1st (well, technically 2nd) command line argument as the number of
1284 frames to look at.
1285 """
1286
1287 if nframes < 1:
1288     # complains if it's asked to look at less than 1 frame (and gives up)
1289     print("How the hell am I supposed to look at less than 1 frame!?")
1290     sys.exit(1)
1291
1292
1293 objectPositions: Dict[str, List[Vector3D]] = {
1294     "cyan": [],
1295     "red": [],
1296     "white": [],
1297     "blue": [],
1298     "green": [],
1299     "yellow": [],
1300     "orange": []
1301 }
1302 """
1303 This is a dictionary which will hold the positions of the objects for every
1304 frame.
1305 """
1306
1307 print("frame  identity  distance") # header for the required frame data info.
1308
1309 # the following 10 lines were adapted from the assignment brief.
1310 for frame in range(0, nframes):
1311     # we work out the filenames for the left and right images for this frame,
1312     # and then we open those images using opencv.
1313     # (and also check to see if the images could actually be opened.)
1314     fn_left: str = sys.argv[2] % frame
1315     im_left: np.ndarray = cv2.imread(fn_left) # left image (BGR)
1316     checkIfImageWasOpened(fn_left, im_left)
1317
1318     fn_right: str = sys.argv[3] % frame
1319     im_right: np.ndarray = cv2.imread(fn_right) # Right image (BGR)
1320     checkIfImageWasOpened(fn_right, im_right)
1321
1322     if debugging:
1323         """
1324         You remember those debugging functions from earlier, right?
1325         Well, this is where they get used. If you enabled 'debugging' ofc.
1326         """
1327         debug(fn_left, fn_right, im_left, im_right, frame)
1328         # END OF DEBUGGING CODE
1329
1330     posXYZ: Dict[str, Vector3D] = \
1331         calculateAndPrintPositionsOfObjects(im_left, im_right, frame)
1332     """
1333     We obtain the identifiers and XYZ positions of all the objects that are
1334     present within both of the stereo frames.
1335     """
1336
1337     for o in [*posXYZ.keys()]:
1338         objectPositions[o].append(posXYZ[o])
1339         # and we append them to the list of all positions for that object.
1340
1341 # Finally, we print out what is/isn't a UFO.
1342 print("UFO:{0}".format(makeUfoString(objectPositions)))
1343
1344 """
1345 That's all, folks!
1346 """

```